

EXPERIENCE

CEO & Co-Founder, FAR.AI.

Jan 2022–present

FAR.AI is an AI safety research non-profit. As CEO, I set the overall agenda for our research projects and manage our technical staff to execute on this agenda.

Director, METR.

May 2024–present

I serve on the board of Model Evaluation & Threat Research (METR) who research, develop and run evaluations to understand the capabilities and risks of frontier AI systems.

Trustee & Director, LISA.

Jul 2024–present

The London Initiative for Safe AI (LISA) cultivates a home for leading AI safety research in London and facilitates mentorship and upskilling programmes.

Director, SAIF.

Dec 2024–present

The Safe AI Forum (SAIF) fosters responsible governance to reduce catastrophic risks through shared understanding and collaboration among key global actors.

Research Intern, DeepMind.

Jan 2021–May 2021; May 2019–Oct 2019

DeepMind is an AI research lab. I worked with Dr Geoffrey Irving and Dr Jan Leike leading to the papers “Uncertainty Estimation for Language Reward Models” and “Quantifying Differences in Reward Functions”.

EDUCATION

University of California, Berkeley, PhD in Artificial Intelligence.

2017–2022

My thesis *Towards Trustworthy Machine Learning* focused on techniques for advanced automated systems to reliably act in accordance with human preferences. I was advised by Prof. Stuart Russell.

University of Cambridge, MPhil in Advanced Computer Science.

2015–2016

Graduated with **distinction**. Awarded **Best Student Prize** (1st out of 31 students).

University of Cambridge, BA (Hons) in Computer Science.

2012–2015

Graduated with **first class** degree. Awarded **Best Student Prize** in 2014, ranking **1st** out of 80 students, and in other years achieved a result in the top 10%.

SELECTED PUBLICATIONS

Mohammad Tafeeque, Stefan Heimersheim, **Adam Gleave**, Chris Cundy. “The Obfuscation Atlas: Mapping Where Honesty Emerges in RLVR with Deception Probes”. In *International Conference on Machine Learning (ICML)*, 2026. **Spotlight Poster**.

Chris Cundy, **Adam Gleave**. “Preference Learning with Lie Detectors can Induce Honesty or Evasion”. In *NeurIPS*, 2025.

Brendan Murphy, Dillon Bowen, ..., **Adam Gleave**, Kellin Pelrine. “Jailbreak-Tuning: Models Efficiently Learn Jailbreak Susceptibility”. In *Empirical Methods in Natural Language Processing*, 2025.

Nikolhaus Howe, Michał Zajac, Ian McKenzie, ..., **Adam Gleave**. “Scaling Trends in Language Model Robustness”. In *International Conference on Machine Learning (ICML)*, 2025. **Spotlight Poster** (top 2.6%).

Tom Tseng, Euan McLean, Kellin Pelrine, Tony T Wang, **Adam Gleave**. “Can Go AIs be adversarially robust?”. In *AAAI Conference on Artificial Intelligence*, 2025.

Dillon Bowen, Brendan Murphy, Will Cai, David Khachaturov, **Adam Gleave**, Kellin Pelrine. “Scaling Trends for Data Poisoning in LLMs”. In *AAAI Conference on Artificial Intelligence*, 2025.

Joar Skalse, ..., **Adam Gleave**, Alessandro Abate. “STARC: A General Framework For Quantifying Differences Between Reward Functions”. In *International Conference on Learning Representations*, 2024.

Tony Tong Wang, **Adam Gleave**, Tom Tseng, Kellin Pelrine, ..., Stuart Russell. “Adversarial Policies Beat

Superhuman Go AIs”. In *International Conference on Machine Learning*, 2023. **Oral Paper** (top 10% of accepted papers); and **Best Paper Award** at NeurIPS ML Safety Workshop 2022.

Joar Skalse, ..., Stuart Russell, **Adam Gleave**. “Invariance in Policy Optimisation and Partial Identifiability in Reward Learning”. In *International Conference on Machine Learning*, 2023.

Adam Gleave, ..., Stuart Russell and Jan Leike. “Quantifying Differences in Reward Functions”. In *International Conference on Learning Representations*, 2021. **Spotlight Paper** (top 20% of accepted papers).

Adam Gleave, ..., Neel Kant, Sergey Levine and Stuart Russell. “Adversarial Policies: Attacking Deep Reinforcement Learning”. In *International Conference on Learning Representations*, 2020.

SELECTED TECHNICAL REPORTS AND WORKSHOP PAPERS

Ann-Kathrin Dombrowski, Dillon Bowen, **Adam Gleave**, Chris Cundy. “The Safety Gap Toolkit: Evaluating Hidden Dangers of Open-Source Models”. In arXiv, 2025.

Ian R McKenzie, Oskar J Hollinsworth, Tom Tseng, Xander Davies, Stephen Casper, Aaron D Tucker, Robert Kirk, **Adam Gleave**. “STACK: Adversarial Attacks on LLM Safeguard Pipelines”. In arXiv, 2025.

Yoshua Bengio, Tegan Maharaj, Luke Ong, Stuart Russell, ..., **Adam Gleave**, et al. “The Singapore Consensus on Global AI Safety Research Priorities”. In arXiv, 2025.

Mohammad Taufeeque, Aaron David Tucker, **Adam Gleave**, Adrià Garriga-Alonso. “Interpreting learned search: finding a transition model and value function in an RNN that plays Sokoban”. In arXiv, 2025.

Matthew Kowal, Jasper Timm, ..., David Rand, **Adam Gleave**, Kellin Pelrine. “It’s the Thought that Counts: Evaluating the Attempts of Frontier LLMs to Persuade on Harmful Topics”. In arXiv, 2025.

Dillon Bowen, Ann-Kathrin Dombrowski, **Adam Gleave**, Chris Cundy. “AI Companies Should Report Pre-and Post-Mitigation Safety Evaluations”. In arXiv, 2025.

Lewis Hammond, Alan Chan, ..., **Adam Gleave**. “Multi-Agent Risks from Advanced AI”. In arXiv, 2025.

Adrià Garriga-Alonso, Mohammad Taufeeque, **Adam Gleave**. “Planning behavior in a recurrent neural network that plays Sokoban”. In *ICML MI Workshop*, 2024.

Kellin Pelrine, Mohammad Taufeeque, Michał Zając, Euan McLean, **Adam Gleave**. “Exploiting novel GPT-4 APIs”. In arXiv, 2023.

Lauro Langosco, David Krueger, **Adam Gleave**. “Training Equilibria in Reinforcement Learning”. *NeurIPS Deep RL Workshop*, 2022.

Lev E McKinney, Yawen Duan, David Krueger, **Adam Gleave**. “On The Fragility of Learned Reward Functions”. *NeurIPS ML Safety Workshop*, 2022.

Erik Jenner, Joar Max Viktor Skalse, **Adam Gleave**. “A general framework for reward function distances”. *NeurIPS ML Safety Workshop*, 2022.

Adam Gleave, Mohammad Taufeeque, Juan Rocamonde, ..., Maximilian Ernestus, Nora Belrose, Scott Emmons, Stuart Russell. “imitation: Clean Imitation Learning Implementations”. In arXiv, 2022.

Pavel Czempin, **Adam Gleave**. “Reducing Exploitability with Population Based Training”. In *New Frontiers in Adversarial Machine Learning at ICML*, 2022.

Adam Gleave, Geoffrey Irving. “Uncertainty Estimation for Language Reward Models”. In arXiv, 2022.

Eric J. Michaud, **Adam Gleave**, Stuart Russell. “Understanding Learned Reward Functions”. In *NeurIPS DeepRL Workshop*, 2020.

Aaron Tucker, **Adam Gleave**, Stuart Russell. “Inverse reinforcement learning for video games”. In *NeurIPS DeepRL Workshop*, 2018.

Adam Gleave, Oliver Habryka. “Multi-task Maximum Causal Entropy Inverse Reinforcement Learning”. In *GoalsRL Workshop at ICML/IJCAI/AAMAS*, 2018.

Sören Mindermann, Rohin Shah, **Adam Gleave**, Dylan Hadfield-Menell. “Active Inverse Reward Design”. In *GoalsRL Workshop at ICML/IJCAI/AAMAS*, 2018.

COMMUNITY CONTRIBUTIONS

International Association for Safe and Ethical AI (IASEAI) (2024–present): Program Chair of IASEAI 2025, Workshop of IASEAI 2026, serving on IASEAI Advisory council since its formation.

Open-source software. Maintainer of [Stable Baselines](#), [Stable Baselines3](#) and [imitation](#), implementations of RL and imitation learning algorithms, with a total of over 10000 stars on GitHub.

Organizer of Center for Human Compatible AI Workshop 2023, 2022; Symposium on AGI Safety 2023, 2022; NeurIPS 2019 Human-Aligned AI Social.

Fund Manager for the Long-Term Future Fund (2020–2022), working with 3-5 other managers to direct over \$6 million in grant funding to effective charities working to improve humanity’s long-term future.

SELECTED PRESS COVERAGE

Can AI be superhuman? Flaws in top gaming bot cast doubt

Nature; *Scientific American*. July 2024.

Man beats machine at Go in human victory over AI

Financial Times. February 2023.

Can Your Career Help Change the World?

Wall Street Journal. October 2021.

Reinforcement-learning AIs are vulnerable to a new kind of attack

MIT Technology Review. February 2020.

SELECTED INVITED TALKS AND INTERVIEWS

Securing Advanced AI: Capabilities, Risks and Solutions

Technical AI Safety Conference, Japan and *UK AISI AI Alignment Conference, UK*. April and October 2025.

ML systems are persistently vulnerable

Guest Lecture for AI Safety Course, IIT Hyderabad. February 2025.

International Cooperation on AI Safety

RAND; *London Initiative for Safe AI*. September 2024.

Towards Model Harmlessness

8th Annual Center for Human-Compatible AI Workshop. June 2024.

ML Systems Are Persistently Vulnerable

Recurse Center and *Google DeepMind*. April and May 2024.

Multi-Agent Vulnerabilities in Superhuman AI

Keynote Talk at NeurIPS Multi-Agent Security: Security as Key to AI Safety Workshop. December 2023.

Automated Testing for Machine Learning Systems

Oak Ridge National Laboratory. August 2023.

Panel: The Future of Generalization: Scale, Safety and Beyond (Moderator)

ICML Workshop on Spurious Correlations, Invariance, and Stability. July 2023.

AWARDS

AI2050 Fellowship Recipient, Schmidt Sciences.

2025

AI2050 is a philanthropic initiative that aims to support exceptional people working on key opportunities and hard problems that are critical for society to benefit from AI.

AI Fellowship Recipient, Open Philanthropy.

2021

Fellowship support for work improving the safety of transformative AI systems.

Winton Capital Best MPhil Student Prize, University of Cambridge.

2016

Awarded for the best result in the MPhil in Advanced Computer Science.